
SUMMARY OF QUALIFICATIONS

- Strong experience in applying ML end-to-end from defining the problem, collecting/crowdsourcing the required data, research on state-of-the-art techniques/models, and to shipping production ready systems consumable by the end user.
 - Experience in NLP, deep learning, chatbots, unstructured data and general ML (classification, regression, unsupervised).
 - Innate ability to deliver meaningful insights and communicate clearly with the cross-functional team, stakeholders, and users.
 - Technical: Python, R, Ruby (beginner), SQL, Hive, Spark, Tableau, ML frameworks (Sklearn, Tensorflow, Keras)
-

EDUCATION

University of Washington, Information School, Seattle, WA (3.93/4.0) June '17

Master of Science in Information Management -- Specialization in Data Science & Analytics. Relevant Coursework: Data Science, Machine Learning & Econometrics, Advanced Data Visualization, Information Architecture, Data Scaling, Applied Social Statistics.

Osmania University, Hyderabad, India (86.85%) June '14

B.S. Computer Science and Engineering -- Relevant Coursework: Data Structures, Software Engineering, Database Management Systems, Artificial Intelligence, Data Mining, Information Security, Information Storage, and Management

WORK EXPERIENCE

SAP Labs, Conversational AI, Palo Alto, CA – Data Scientist / ML Engineer March '17 – Present

- End to end applied research, development, and deployment of machine learning models in the NLP domain.
- Projects include natural language understanding models for intent recognition, entity extraction, coreference resolution, closed domain question & answering on structured/unstructured text, and information retrieval systems.
- Communicating and collaborating with the product team, business stakeholders and engineers to ship features on the platform.
- Shipped features in production: Demo link - [pronoun resolution for dialog context management, text augmentation as a service](#).

Concur, Seattle, WA – Product Data Science Intern June '16 – March '17

- Worked in the Product Intelligence team to build models for customer churn to proactively identify and mitigate attrition risks.
- Informed product strategy for the development of a new simplified Analytics Product for Concur Customers to replace an existing 150M ARR reporting product using NLP/ML.
- Built and automated end-to-end data pipeline for data extraction, preprocessing, mining, prediction and data visualization.

Google Inc., Hyderabad, India – Associate Account Strategist May '14 – Aug '15

- Provided strategic advice and data-driven insights to help advertisers maximize ROI by working closely in a consultative role.
- Worked in the Tools team to develop advanced analytics dashboards for the Business Unit to drive business decision making.

Google Inc., Hyderabad, India – SMB Intern May '13 – July '13

- Impact measurement of AdWords campaign optimizations performed by account optimizers.
 - Consolidated and analyzed market-wide data to suggest new optimization opportunities and process improvements.
-

ACADEMIC RESEARCH PROJECTS

Data for Economic Development - Improving access to Financial Institutions based on Mobile Cellular Data in Ghana (2016)

- Worked as a Research Assistant for Prof. Joshua Blumenstock to crunch terabyte scale Call Detail Records data using Spark.
- Cleaning, anonymizing and visualizing data for enabling analysis of the socio-economic behavior of the population.

Healthcare Cost and Utilization Project (HCUP) - Fixed Effects Econometrics (2016) ([Project Link](#))

- Worked with Prof. Benjamin Althouse on modeling relationships between vaccination uptake and hospitalization for WA state.
- Built fixed effects model in R for gauging the impact of vaccination within counties normalizing for demographic characteristics.

Wikitrends - Wikipedia Trend Analysis During US Presidential Elections (2016) ([Project Link](#))

- Understanding relationships between Wikipedia and general interest on Google Trends over time for presidential elections.
- Mined extremely large messy data in the form of Wikipedia log files using Turi (Apple) Dato Distributed API on 20 node cluster.